

Yunbei Zhang

✉ yunbeizhang.ml@gmail.com ☎ +1-408-896-0102 📍 New Orleans, LA 70018

💻 yunbeizhang.github.io 🎓 Scholar [in](#) LinkedIn [GitHub](#)

RESEARCH INTERESTS

- **Test/Inference-Time Learning:** Improving reasoning via RL with unverified rewards, and adapting to new and changing domains.
- **Safety & Alignment:** Ensuring frontier agents are robust, reliable, and aligned with human values.
- **Efficient Model Re-use:** Lightweight reprogramming and prompting for adapting pre-trained models.

EDUCATION

- **Tulane University** Aug 2022 - Present
Ph.D. in Computer Science New Orleans, LA, USA
 - Focus: Developing efficient, robust, and generalizable ML methods for real-world deployment.
- **Osaka University** Apr 2018 - Mar 2022
B.S. in Mathematics (Award of Excellence) Osaka, Japan
 - Honors: Osaka University Award of Excellence, Multiple Merit-based Scholarships
- **Peking University** Jul 2013 - May 2015
Incomplete Beijing, China
 - Withdrew after two years to pursue alternative academic and career opportunities.

PROFESSIONAL EXPERIENCE

- **Waymo LLC** Jun 2026 - Aug 2026 (Expected)
Research Scientist Intern Mountain View, CA, USA
 - Multi-modal LLM post-training, Prediction and Planning
- **Oak Ridge National Laboratory** Jan 2026 - April 2026
Research Intern Oak Ridge, TN, USA
 - Uncertainty Quantification for Large Foundation Model and Multi-modal LLMs
- **Amazon** Sep 2025 - Dec 2025
Applied Scientist Intern Seattle, WA, USA
 - Multi-modal LLM RL-based post-training and Safety
 - Automatic Multimodal Red Teaming via Agentic Planning
- **KLA Corporation** May 2025 - Aug 2025
Machine Learning Research Intern Milpitas, CA, USA
 - Applied VFM Knowledge Distillation for compact Object Detection
- **Tulane University** Jan 2024 - May 2024
Teaching Assistant (Outstanding TA Award 2024) New Orleans, LA, USA
 - Received Outstanding Teaching Assistant Award for exceptional student mentorship

SELECTED PUBLICATIONS (*equal contribution)

Conference Papers:

- [CVPR'26] Prime Once, then Reprogram Locally: An Efficient Alternative to Black-Box Service Model Adaptation (**Highlight**)
Y. Zhang, C. Cai, F. Liu, J. Hamm
- [CVPR'26] Seeing Clearly, Reasoning Confidently: Plug-and-Play Remedies for Vision Language Model Blindness
X. Hu, H. Ni, Y. Zhang, J. Hamm, Z. Li, Z. Ding
- [ACL'26 Main] Not All Directions Matter: Toward Structured and Task-Aware Low-Rank Adaptation
X. Xiao*, C. Ma*, Y. Zhang*, C. Liu, Z. Wang, Y. Li, L. Zhao, G. Hu, T. Wang, H. Xu
- [ICML'25] DPCore: Dynamic Prompt Coreset for Continual Test-Time Adaptation
Y. Zhang, A. Mehra, S. Niu, J. Hamm

- [NeurIPS'25] Doctor Approved: Generating Medically Accurate Skin Disease Images through AI-Expert Feedback
J. Wang*, **Y. Zhang***, Z. Ding, J. Hamm
- [NeurIPS'24] Understanding the Transferability of Representations via Task-Relatedness
A. Mehra, **Y. Zhang**, J. Hamm
- [TMLR] Prompt-based Adaptation in Large-scale Vision Models: A Survey
Y. Zhang*, X. Xiao*, L. Zhao*, et al.
- [ICASSP'26] CTR-LoRA: Curvature-Aware and Trust-Region Guided Low-Rank Adaptation for Large Language Models (**Oral**)
Z. Wang, M. Mo, X. Xiao, C. Liu, C. Ma, **Y. Zhang**, et al.
- [USENIX'25] SoK: Can Synthetic Images Replace Real Data? A Survey of Utility and Privacy
Y. Cheung, **Y. Zhang**, N. Marrouche, J. Hamm
- [WACV'25] OT-VP: Optimal Transport-guided Visual Prompting for Test-Time Adaptation (**Oral**)
Y. Zhang, A. Mehra, J. Hamm
- [ACM MM'25] Visual Instance-aware Prompt Tuning
Y. Zhang*, X. Xiao*, X. Li, T. Wang, X. Wang, Y. Wei, J. Hamm, M. Xu
- [ACM MM'25] eSkinHealth: A Multimodal Dataset for Neglected Tropical Skin Diseases
J. Wang, X. Hu, **Y. Zhang**, et al.
- [WACV'26] RoadBench: A Vision-Language Foundation Model for Road Damage Understanding
X. Xiao, **Y. Zhang**, J. Wang, et al.
- [WACV'24] On the Fly Neural Style Smoothing for Risk-Averse Domain Generalization
A. Mehra, **Y. Zhang**, B. Kailkhura, J. Hamm

Workshop Papers:

- [ICLRW'26] Adapting in the Dark: Towards Stable and Efficient Black-Box Test-Time Adaptation (**Oral**)
- [ICLRW'26] Visual Exclusivity Attacks: Automatic Multimodal Red Teaming via Agentic Planning (**Oral**)
- [ICLRW'26] Agents in the Wild: Safety, Society, and the Illusion of Sociality on Moltbook
- [ICMLW'25] Describe Anything in Medical Images
- [CVPRW'24] Achieving Reliable and Fair Skin Lesion Diagnosis via Unsupervised Domain Adaptation (**Oral**)
- [CVPRW'24] Test-time Assessment of a Model's Performance on Unseen Domains via Optimal Transport

Recent Submissions (Under Review):

- Wan-R1: Verifiable-Reinforcement Learning for **Video Generation**
- Are **Medical Vision-Language Foundation Models** Ready for Dermatology
- ClawSafety: "Safe" LLMs, Unsafe Agents
- Continual Test-Time Adaptation: A Comprehensive Survey

HONORS & AWARDS

• Tinker Research Grant <i>Thinking Machines Lab</i>	2026
• Lambda's Research Grant <i>Lambda</i>	2025
• Outstanding Teaching Assistant Award <i>Tulane University</i>	2024
• Osaka University Award of Excellence <i>Osaka University</i>	2019
• Monbukagakusho Honors Scholarship <i>Ministry of Education, Japan</i>	2018, 2021
• School of Science Honor Scholarship (4 years) <i>Osaka University</i>	2018-2021
• Ichikawa International Scholarship <i>Osaka University</i>	2021

ACADEMIC SERVICE

- **Conference Reviewer:** NeurIPS 2025 (**Top Reviewer**), ICLR 2025-2026, ACM MM 2025-2026, CVPR 2026, ECCV 2026, ICML 2026.
- **Journal Reviewer:** Transactions on Machine Learning Research
- **Volunteer:** NeurIPS 2023 Outreach Program

TECHNICAL SKILLS

- **Languages:** Python, C++, MATLAB, Bash, $\text{\LaTeX}$
- **ML Frameworks:** PyTorch, TensorFlow, JAX, Hugging Face Transformers, scikit-learn
- **Tools:** Git, Docker, CUDA, Weights & Biases, AWS, Google Cloud Platform
- **Natural Languages:** English (Fluent), Japanese (Fluent), Chinese (Native)